

Interpreting Agent Behavior (IAB)

Human-Centered Interpretation of Agents

COLM 2026 Workshop Proposal · [Proposal Website](#)

1. Description

Scope. The widely deployed commercial agents such as Claude and Codex have exhibited complex behaviors, including planning, reasoning, tool use, and failure recovery [38, 33]. These behaviors are captured in vast amounts of behavioral data such as execution logs, prompting histories, and interaction traces [25]. However, existing approaches to understanding agent behaviors remain largely outcome-focused [28]: benchmarks can tell us whether an agent succeeds or fails [18, 19, 20, 41, 36, 26], they cannot tell us *why* [12]. For example, a fail rate does not reveal that an agent fails because it cannot select effective tools [24]. Understanding the “why” is what humans actually need: developers need it to debug agents [22], researchers need it to identify behavioral patterns [25], and policymakers need it to ensure agents align with human values and expectations [28]. **In this workshop, we are interested in a fundamental question: how can we help humans understand agent behaviors?** Humans cannot read thousands of log entries. They need patterns, summaries, and explanations, in other words, *interpretation*. Building human-understandable levels of interpretation is essential for evaluating agents and designing agents align with human expectations and benefit human life.

Importance and timeliness. IAB is important and timely because real-world agent deployments have reached unprecedented scale by early 2026. Agents behavioral data are becoming increasingly diverse, from execution traces [27, 16], tool-use logs [32, 24], decision records [34, 37], and simulated behaviors [13, 30]. Humans need to understand these behaviors, for example, why certain phenomena such as failures occur [12]. Researchers and practitioners are actively adopting and developing empirical methods for understanding agent behaviors, including grounded theory [12, 29], qualitative analysis [21, 23], error analysis [12, 14], corpus analysis [35, 40, 39], trace analysis [27, 11, 15, 10, 17], and red-teaming [31].

Building community. IAB aims to establish a new community that brings together researchers from interpretability, evaluation, alignment, computational social science, behavioral science, collective intelligence, responsible AI, governance, and HCI. Each community contributes unique methodological perspectives and can benefit from the questions IAB addresses. To grow the community beyond the workshop, we will (i) keep a proceeding of accepted papers with **recordings** of keynotes, talks, and the panel; and (ii) a **repository** of behavioral datasets, annotation schemas, and analysis tools contributed by authors and organizers where available. The organizers are also already collaborating on a **community-authored paper** synthesizing shared vocabulary, open problems, and methodological guidelines informed by accepted submissions and the panel discussions.

Supporting junior researchers and underrepresented groups. We accept non-archival submissions to lower the barrier for early-stage work to receive timely feedback. During the workshop, we reserve spotlight slots for student-led papers and poster sessions to facilitate direct exchange between junior and senior researchers. Keynote speakers and panelists will be encouraged to share mentorship advice. To ensure that junior researchers’ work receives recognition, we offer best paper awards for student-led work. Moreover, we have applied for institutional funding (e.g., from JHU and Google) to provide travel grants and registration waivers for underrepresented students, and we support remote attendance for those who cannot travel.

Expected attendance. We anticipate 30-50 submissions and 60–100 attendees. At COLM 2025, the closely related INTERPLAY and ORIGen workshops accepted 25 and 12 papers [2, 4], and agent-focused workshops such as *AgentWild* (ICLR 2026) [5] and *MTI-LLM* (NeurIPS 2025) [3] have drawn even larger audiences. Meanwhile, the growing popularity of agents has driven increasing work on behavior analysis [27, 12, 14]. In addition, our confirmed keynote speakers and panelists

from top institutions (MIT, Stanford, OpenAI, and CMU) are expected to further attract a broad audience. Together with organizers and advisory board members across 9+ institutions, we are well positioned to reach diverse communities.

2. Submission Format and Reviewing

We solicit two types of non-archival submissions: (1) long papers (up to 9 pages + references) and (2) short papers (up to 4 pages + references) according to the COLM template [7].

We welcome four key contributions, including: **1) Datasets and Resources** (e.g., annotated interaction traces, failure catalogs, human-agent conversational data, simulated behavior data); **2) New Methods and Tools** (e.g., qualitative coding frameworks, think-aloud protocols, interaction visualization, credit assignment tools, LLM-as-judge pipelines); **3) Empirical Studies** (e.g., error taxonomies, failure mode analysis, case studies of agent decisions, probing studies of internal states); **4) Human-Agent Interaction Analysis** (e.g., user verification strategies, user intent grounding, mid-execution adjustment, collaborative failure recovery).

We also encourage negative results and methodological position papers. Submissions will use COLM style files and be double-blind, with each paper receiving at least two reviews. We will use OpenReview for submission management. To ensure reviewing rigor, we will form a **Program Committee** beyond the organizing committee: the 9 organizers will serve as primary reviewers and area chairs, and we will additionally recruit 15–20 external reviewers from our advisory board’s networks and prior workshop PCs to keep per-reviewer load manageable and to handle conflicts of interest [8]. Conflicts of interest will follow the COLM COI policy.

Timeline: Submission deadline: June 23, 2026 · Notification: July 24, 2026 · Workshop: Oct 09, 2026.

3. Workshop Schedule

- | | |
|---|--|
| • 09:00–09:10 Opening Remarks | • 13:10–13:45 Keynote 3: Bowen Baker |
| • 09:10–09:45 Keynote 1: Armando Solar-Lezama | • 13:45–15:00 Paper Presentations (5 × 15 mins) |
| • 09:45–10:20 Keynote 2: Diyi Yang | • 15:00–15:50 Poster Session #2 + Coffee Break |
| • 10:20–10:50 Paper Presentations (2 × 15 mins) | • 15:50–16:35 Panel (Armando Solar-Lezama, Diyi Yang, Bowen Baker, Graham Neubig): <i>Empirical Methods for Understanding Agent Behavior</i> |
| • 10:50–11:40 Poster Session #1 + Coffee Break | • 16:35–16:50 Best Paper Award + Closing Remarks |
| • 11:40–12:10 Paper Presentations (2 × 15 mins) | |
| • 12:10–13:10 Lunch | |

Invited keynote speakers (all confirmed): **Armando Solar-Lezama** (MIT CSAIL) – Professor; pioneered sketch-based program synthesis and neurosymbolic programming. **Diyi Yang** (Stanford) – Assistant Professor; research on human-centered NLP, social dynamics, and human-AI interaction. **Bowen Baker** (OpenAI) – Research Scientist; research on multi-agent systems and emergent behavior. **Panel discussion:** **Armando Solar-Lezama**, **Diyi Yang**, **Bowen Baker**, and **Graham Neubig** (CMU LTI – Associate Professor; research on language agents and code intelligence; creator of OpenHands). We also plan to invite more panelists to join the discussion.

4. Related Workshops

IAB is the first workshop on empirical methods for building human-understandable interpretations of agent behavior in real-world scenarios. First, IAB’s emphasis on *behavior interpretation* distinguishes it from workshops centered on agent development, improvement, and model internals, including *AgentWild* (ICLR 2026) [5], *MTI-LLM* (NeurIPS 2025) [3], *INTERPLAY* (COLM 2025) [2], and *BlackboxNLP* (EMNLP 2025) [1]. Second, workshops on human-centered evaluation such as *ORIGen* (COLM 2025) [4], *HEAL* (CHI 2026) [9], and *Human-Agent Collaboration* (CHI 2026) [6], those workshops primarily focus on evaluating and auditing the quality of model outputs. IAB is more narrowly focused on interpreting observable agent execution behavior in real-world scenarios, including how they plan, use tools, fail, and recover, a problem that has mainly emerged with the recent widespread deployment of agentic systems.

5. Organizer Bios

Our 9-member committee from 5 institutions ensures diversity across multiple dimensions: **Career stage:** 1 assistant professor, 4 postdoctoral researchers, 1 PhD student, and 3 industry research scientists. **Sector:** both representation in academia (6) and industry (3). **Research background:** a mix of AI/NLP researchers and HCI researchers. **Geography:** members span institutions across the US and Singapore. **Gender:** the committee includes 4 women, 4 men and 1 non-binary individual.

Jie (Sophia) Gao [Academia] [HCI] (jgao77@jh.edu) is a Malone Postdoctoral Fellow at Johns Hopkins University. Her research focuses on human-AI collaboration for subjective tasks, including AI-assisted qualitative data analysis and developer interactions with coding agents. Her paper at TOCHI2023 and CHI 2024, which pioneered for AI-assisted qualitative analysis. She has also organized the "LLMs as Research Tools" workshop at CHI2024. *Contribution: Jie brings 5+ years of expertise in AI-assisted qualitative analysis method and HCI research methods to agent behavior analysis.*

Kaiser Sun [Academia] [AI] (hsun74@jhu.edu) is a PhD student at the Center for Language and Speech Processing and Data Science and AI Institute at Johns Hopkins University, advised by Mark Dredze. Their research focuses on the training dynamics, interpretability, and multimodality of LLMs. Their work has been recognized with a Best Paper Honorable Mention at CoNLL 2023 and coverage by MIT Technology Review. Previously, they gained industry experience at Meta AI and AWS AI Labs. *Contribution: Kaiser brings expertise in LLM interpretability and training dynamics to agent behavior analysis.*

Teresa Yeo [Industry] [AI] (teresa.ysq@gmail.com) is a Research Scientist at Google DeepMind. Her research focuses on understanding and steering foundation model behavior under distributional and task shifts, with particular interest in how models can be monitored and adapted when their behavior departs from expectations, a critical challenge as models become agentic. She co-organized the 2nd and 3rd Test-time Adaptation Workshop (ICML'25 and ICLR'26), the 1st Workshop on Monitoring ML Models Under Drift (ICLR'26), and was an invited speaker at the 1st Test-time Adaptation Workshop (CVPR'24). *Contribution: Teresa brings expertise in model behavior monitoring and adaptation to the workshop, with a focus on real-world deployment challenges.*

Daniel Khashabi [Academia] [AI] (danielk@jhu.edu) is an assistant professor in the Department of Computer Science at Johns Hopkins University and is affiliated with the Data Science and AI Institute and the Center for Language and Speech Processing (CLSP). His research is anchored around natural language as a key medium for communication, driven by the goal of making these systems more helpful, reliable, and efficient. He is the recipient of AI2's Lasting Impact Award (2024), multiple industry faculty awards (Amazon, Google, Apple), and several Best Paper Awards (COLM 2024, ACL 2023, NAACL 2022). His teaching has awarded him the Joel Dean Award for Excellence in Teaching in the Department of Computer Science in 2025. *Contribution: Daniel brings expertise in language-driven reasoning systems and their evaluation to the workshop, with a focus on reliability and real-world impact.*

Zhuoran Lu [Industry] [HCI] (lu800@purdue.edu) is a Senior Research Scientist at Accenture. His research lies at the intersection of HCI and AI, with a focus on human-AI interaction, human-centered AI, and agentic human-AI collaboration. His work has been published in top-tier venues across HCI (e.g., CHI, CSCW, IUI) and AI (e.g., ICLR, EMNLP, AACL), and was recognized with the Best Paper Award at CSCW 2022. *Contribution: Zhuoran brings expertise in human-AI interaction and human-centered AI design to the understanding of human-agent collaboration.*

Boyuan Zheng [Industry] [AI] (steven.zheng010@gmail.com) is a Member of Technical Staff at xAI. His research focuses on building language agents for digital environments, including scaling RL training for computer-use agents, self-play, and automated research. He's a core contributor to the computer-use and web dev capabilities of Grok 4.20 and has published widely on computer-use

agents, including MolmoWeb, SeeAct, UGround, and Mind2Web. He co-organized the Workshop on Computer Use Agents at ICML 2025. *Contribution: Boyuan brings expertise in building and training language agents for digital environments, providing insights into real-world agent behavior from an industry deployment perspective.*

Katherine Van Koeving [Academia] [HCI] (kvankoe1@jh.edu) is a postdoctoral fellow at the Johns Hopkins Data Science and AI Institute. She received her PhD from Cornell University. Her research focuses on the flow of people and ideas through networks and hidden linguistic gradients in digital communication. *Contribution: Katherine brings expertise in computational social science, computational linguistics, and digital communication analysis to the workshop, providing insights into how information propagates through human-agent interactions.*

Sijie Ji [Academia] [AI] (sijieji@caltech.edu) is a Schmidt Science Fellow at Caltech. She develops engineering principles that expand the observability of real-world systems, allowing AI to reason about human behavior, physiology, and hidden internal states beyond traditional sensing methods. Dr. Ji has served as an organizer and program committee member for conferences such as SenSys, BuildSys, MobiSys, MobiCom, and CPS-IoT Week. She has been recognized as an NSF CPS Rising Star, an NIH mHealth Training Institute Scholar, and a recipient of the N2Women Young Researcher Fellowship. *Contribution: Sijie brings expertise in cyber-physical systems, real-world observability, and AIoT deployment, contributing a systems perspective on how agent behavior should be interpreted under real-world constraints, noisy interaction environments, and human-in-the-loop workflows.*

Jen-tse Huang [Academia] [AI] (jhuan236@jh.edu) is a postdoctoral researcher at Johns Hopkins University CLSP. His work focuses on LLM evaluation and cognitive alignment, with publications at NeurIPS, ICLR, and ICML. He organized the PersonaLLM Workshop at NeurIPS 2025. *Contribution: Jen-tse brings expertise in multi-agent simulation and human-AI cognitive alignment, providing insights into how to design evaluations that capture the nuances of agent behavior and human interpretation.*

Advisory Board: We also have an advisory board of 7 senior researchers who will provide feedback on the workshop’s direction, help identify keynote speakers and panelists, and promote the workshop within their networks. The advisory board further strengthens the committee’s diversity, spanning academia and industry, multiple institutions, and both AI and HCI backgrounds: The advisory board includes: **Ziang Xiao** (JHU) · **Soufiane Hayou** (JHU) · **Toby Jia-Jun Li** (Notre Dame) · **Hang Jiang** (Northeastern) · **Weiyan Shi** (Northeastern) · Samuel Nathanson (xAI) · **Wei Lu** (NTU).

References

- [1] BlackboxNLP 2025: Analyzing and interpreting neural networks for NLP. <https://blackboxnlp.github.io/2025/>, 2025.
- [2] INTERPLAY workshop at COLM 2025. <https://interplay-workshop.github.io/>, 2025.
- [3] MTI-LLM: Multi-turn interactions of LLMs workshop at NeurIPS 2025. <https://workshop-multi-turn-interaction.github.io/>, 2025.
- [4] ORIGen workshop at COLM 2025. <https://origen-workshop.github.io/2025/programme/>, 2025.
- [5] AgentWild: Workshop on agent safety and security at ICLR 2026. <https://agentwild-workshop.github.io/>, 2026.
- [6] Human-agent collaboration workshop at CHI 2026. <https://chi26workshop-human-agent-collaboration.hailab.io/>, 2026.

- [7] COLM 2026 call for papers. <https://colmweb.org/cfp.html>, 2026.
- [8] COLM 2026 conflict of interest policy. <https://colmweb.org/coi-policy.html>, 2026.
- [9] HEAL: Human-centered evaluation of language models workshop at CHI 2026. <https://heal-workshop.github.io/>, 2026.
- [10] Gagan Bansal, Jennifer Wortman Vaughan, Saleema Amershi, Eric Horvitz, Adam Fourney, Hussein Mozannar, Victor Dibia, and Daniel S. Weld. Challenges in human-agent communication, 2024. URL <https://arxiv.org/abs/2412.10380>.
- [11] Islem Bouzenia and Michael Pradel. Understanding software engineering agents: A study of thought-action-result trajectories, 2025. URL <https://arxiv.org/abs/2506.18824>.
- [12] Mert Cemri, Melissa Z. Pan, Shuyi Yang, Lakshya A. Agrawal, Bhavya Chopra, Rishabh Tiwari, Kurt Keutzer, Aditya Parameswaran, Dan Klein, Kannan Ramchandran, Matei Zaharia, Joseph E. Gonzalez, and Ion Stoica. Why do multi-agent llm systems fail?, 2025. URL <https://arxiv.org/abs/2503.13657>.
- [13] Jiawei Chen, Ruoxi Xu, Boxi Cao, Ruotong Pan, Yunfei Zhang, Yifei Hu, Yong Du, Tingting Gao, Yaojie Lu, Yingfei Sun, Xianpei Han, Le Sun, Xiangyu Wu, and Hongyu Lin. Towards real-world human behavior simulation: Benchmarking large language models on long-horizon, cross-scenario, heterogeneous behavior traces, 2026. URL <https://arxiv.org/abs/2604.08362>.
- [14] Penny Chong, Harshavardhan Abichandani, Jiyuan Shen, Atin Ghosh, Min Pyae Moe, Yifan Mai, and Daniel Dahlmeier. Talk, evaluate, diagnose: User-aware agent evaluation with automated error analysis, 2026. URL <https://arxiv.org/abs/2603.15483>.
- [15] Peter Cihon, Merlin Stein, Gagan Bansal, Sam Manning, and Kevin Xu. Measuring ai agent autonomy: Towards a scalable approach with code inspection, 2025. URL <https://arxiv.org/abs/2502.15212>.
- [16] Christopher Zhang Cui, Xingdi Yuan, Ziang Xiao, Prithviraj Ammanabrolu, and Marc-Alexandre Côté. Tales: Text adventure learning environment suite, 2025. URL <https://arxiv.org/abs/2504.14128>.
- [17] Liming Dong, Qinghua Lu, and Liming Zhu. Agentops: Enabling observability of llm agents, 2024. URL <https://arxiv.org/abs/2411.05285>.
- [18] C. E. Jimenez et al. SWE-bench: Can language models resolve real-world GitHub issues? In *ICLR*, 2024.
- [19] X. Liu et al. AgentBench: Evaluating LLMs as agents. In *ICLR*, 2024.
- [20] Y. Yao et al. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. In *NeurIPS*, 2024.
- [21] Jie Gao, Yuchen Guo, Gionnieve Lim, Tianqin Zhang, Zheng Zhang, Toby Jia-Jun Li, and Simon Tangi Perrault. Collabcoder: A lower-barrier, rigorous workflow for inductive collaborative qualitative analysis with large language models. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*, CHI ’24, New York, NY, USA, 2024. Association for Computing Machinery. ISBN 9798400703300. doi: 10.1145/3613904.3642002. URL <https://doi.org/10.1145/3613904.3642002>.

- [22] Melody Y. Guan, Miles Wang, Micah Carroll, Zehao Dou, Annie Y. Wei, Marcus Williams, Benjamin Arnav, Joost Huizinga, Ian Kivlichan, Mia Glaese, Jakub Pachocki, and Bowen Baker. Monitoring monitorability, 2025. URL <https://arxiv.org/abs/2512.18311>.
- [23] Jialun Aaron Jiang, Kandrea Wade, Casey Fiesler, and Jed R. Brubaker. Supporting serendipity: Opportunities and challenges for human-ai collaboration in qualitative analysis. *Proc. ACM Hum.-Comput. Interact.*, 5(CSCW1), April 2021. doi: 10.1145/3449168. URL <https://doi.org/10.1145/3449168>.
- [24] Minghao Li, Yingxiu Zhao, Bowen Yu, Feifan Song, Hangyu Li, Haiyang Yu, Zhoujun Li, Fei Huang, and Yongbin Li. API-bank: A comprehensive benchmark for tool-augmented LLMs. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 3102–3116, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.emnlp-main.187. URL <https://aclanthology.org/2023.emnlp-main.187/>.
- [25] Jiaying Lu, Bo Pan, Jieyi Chen, Yingchaojie Feng, Jingyuan Hu, Yuchen Peng, and Wei Chen. Agentlens: Visual analysis for agent behaviors in llm-based autonomous systems. *IEEE Transactions on Visualization and Computer Graphics*, 2024.
- [26] Chang Ma, Junlei Zhang, Zhihao Zhu, Cheng Yang, Yujiu Yang, Yaohui Jin, Zhenzhong Lan, Lingpeng Kong, and Junxian He. Agentboard: An analytical evaluation board of multi-turn llm agents, 2024. URL <https://arxiv.org/abs/2401.13178>.
- [27] Oorja Majgaonkar, Zhiwei Fei, Xiang Li, Federica Sarro, and He Ye. Understanding code agent behaviour: An empirical study of success and failure trajectories, 2025. URL <https://arxiv.org/abs/2511.00197>.
- [28] Mahmoud Mohammadi, Yipeng Li, Jane Lo, and Wendy Yip. Evaluation and benchmarking of llm agents: A survey. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*, KDD '25, page 6129–6139. ACM, August 2025. doi: 10.1145/3711896.3736570. URL <http://dx.doi.org/10.1145/3711896.3736570>.
- [29] Laura K Nelson. Computational grounded theory: A methodological framework. *Sociological methods & research*, 49(1):3–42, 2020.
- [30] Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior, 2023. URL <https://arxiv.org/abs/2304.03442>.
- [31] Ethan Perez, Saffron Huang, Francis Song, Trevor Cai, Roman Ring, John Aslanides, Amelia Glaese, Nat McAleese, and Geoffrey Irving. Red teaming language models with language models. In Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang, editors, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 3419–3448, Abu Dhabi, United Arab Emirates, December 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.emnlp-main.225. URL <https://aclanthology.org/2022.emnlp-main.225/>.
- [32] Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, Sihan Zhao, Lauren Hong, Runchu Tian, Ruobing Xie, Jie Zhou, Mark Gerstein, Dahai Li, Zhiyuan Liu, and Maosong Sun. Toolllm: Facilitating large language models to master 16000+ real-world apis, 2023. URL <https://arxiv.org/abs/2307.16789>.

- [33] Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools, 2023. URL <https://arxiv.org/abs/2302.04761>.
- [34] Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning, 2023. URL <https://arxiv.org/abs/2303.11366>.
- [35] Alex Tamkin, Miles McCain, Kunal Handa, Esin Durmus, Liane Lovitt, Ankur Rathi, Saffron Huang, Alfred Mountfield, Jerry Hong, Stuart Ritchie, Michael Stern, Brian Clarke, Landon Goldberg, Theodore R. Sumers, Jared Mueller, William McEachen, Wes Mitchell, Shan Carter, Jack Clark, Jared Kaplan, and Deep Ganguli. Clío: Privacy-preserving insights into real-world ai use, 2024. URL <https://arxiv.org/abs/2412.13678>.
- [36] Shunyu Yao, Howard Chen, John Yang, and Karthik Narasimhan. Webshop: Towards scalable real-world web interaction with grounded language agents, 2023. URL <https://arxiv.org/abs/2207.01206>.
- [37] Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L. Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models, 2023. URL <https://arxiv.org/abs/2305.10601>.
- [38] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models, 2023. URL <https://arxiv.org/abs/2210.03629>.
- [39] Renwen Zhang, Han Li, Han Meng, Jinyuan Zhan, Hongyuan Gan, and Yi-Chieh Lee. The dark side of ai companionship: A taxonomy of harmful algorithmic behaviors in human-ai relationships. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, CHI '25, New York, NY, USA, 2025. Association for Computing Machinery. ISBN 9798400713941. doi: 10.1145/3706598.3713429. URL <https://doi.org/10.1145/3706598.3713429>.
- [40] Wenting Zhao, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, and Yuntian Deng. Wildchat: 1m chatgpt interaction logs in the wild, 2024. URL <https://arxiv.org/abs/2405.01470>.
- [41] Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. Webarena: A realistic web environment for building autonomous agents, 2024. URL <https://arxiv.org/abs/2307.13854>.